

Actividad: Problemas de regresión

Alan Vilchis A01640260


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score
from sklearn.model_selection import KFold, ShuffleSplit
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import Lasso
from sklearn import linear_model
from sklearn.datasets import load_diabetes
from sklearn.preprocessing import StandardScaler

df = pd.read_csv('/home/alanv/Documents/7/omar/crime_data.csv')
df.drop(['State', 'MR', 'S'], inplace=True, axis=1)
# Dependiente es VR
df.head(5)
```

1. Grafica cada variable predictora vs la variable de respuesta

```
fig, ax = plt.subplots(1,4 , figsize=(15, 5))
ax[0].scatter(df['M'],df['VR'],label='Graph 1')
ax[0].set_title('M vs VR')
ax[1].scatter(df['W'],df['VR'],label='Graph 2')
ax[1].set_title('W vs VR')
ax[2].scatter(df['H'],df['VR'],label='Graph 3')
ax[2].set_title('H vs VR')
ax[3].scatter(df['P'],df['VR'],label='Graph 4')
ax[3].set_title('P vs VR')

plt.show()
```



1. Implementa la fórmula directa para calcular los coeficientes de un modelo de regresión lineal

```
def fit_model(X, y):
    return np.linalg.inv(X.transpose() @ X) @ X.transpose() @ y
# Predict function (it evaluates an array of observations using the
# specified linear model).
def predict(X, beta):
    return X @ beta
```

```
x = np.array([df['M'],df['W'],df['H'],df['P']]).T
X = np.column_stack((np.ones(x.shape[0]), x))
y = np.array(df['VR'])
beta = fit_model(X,y)
print("Coefficients:", beta)
y_pred = predict(X,beta)
y_pred
```

```
Coefficients: [-2014.12471071    9.09424116   -13.08150065
 29.02038939
 63.28190064]
```

```
array([ 471.31673749,  679.90596672,  497.62129419,  855.58247246,
        1191.00967082,  595.55708654,  528.25160529,  593.85344558,
        1024.86921462,  587.23643591,  960.97012249,   96.86685555,
        135.63934346,  762.17751401,  417.81962277,  481.66300289,
        391.73215308, 1448.1420549 ,  667.76406741,  817.53843833,
        -14.06509507,  854.2035039 ,  511.79527333,  624.51787943,
        866.08348497,  286.34197652,  547.78125552,   66.53951321,
        238.12980892,  256.05297423,  753.95073572,  636.29040154,
        529.87014319, 1018.47365009,  600.11397353,  877.44625022,
        515.16447234,  599.88200756,  423.97462462,  888.7265147 ,
        238.47120303,  706.00054226,  829.12156314,  597.30419112,
        478.26299712, -78.13339479,  768.83077574,  478.74145614,
        426.46995063,  251.79993674, 2271.34032589])
```

1. Evalúa con validación cruzada de k-pliegues tu modelo, calculando los valores de R2, MSE y MAE.


```
x = df.iloc[:, 1:].values
y = np.array(df['VR'])
# Evaluate linear regression model using k-fold cross-validation
from sklearn import linear_model
n_folds = 5
kf = KFold(n_splits=n_folds, shuffle = True)
mse_cv = []
mae_cv = []
r2_cv = []
for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]
    #print('a', x_train)
    regr_cv = linear_model.LinearRegression()
    regr_cv.fit(x_train, y_train)
    # Test phase
    x_test = x[test_index]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)
    # Calculate MSE, MAE and R^2
    mse_i = mean_squared_error(y_test, y_pred)
    print('mse = ', mse_i)
    mse_cv.append(mse_i)
    mae_i = mean_absolute_error(y_test, y_pred)
    print('mae = ', mae_i)
    mae_cv.append(mae_i)
    r2_i = r2_score(y_test, y_pred)
    print('r^2= ', r2_i)
    r2_cv.append(r2_i)
print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), ' R^2:',
      np.average(r2_cv))

mse = 188275.9226900601
mae = 298.5080030039882
r^2= 0.6276831267422318
mse = 42428.301883003944
mae = 181.14683877515648
r^2= 0.6319660083367942
mse = 193573.34696674283
mae = 370.72950716622836
r^2= -1.1032376279160534
mse = 9964.174111154396
mae = 89.78830359677059
r^2= 0.8951331197672258
mse = 13641.912032116801
mae = 101.55196814509368
```

```
r^2= 0.7686202057818388
MSE: 89576.73153661561 MAE: 208.34492413744746 R^2:
0.3640329665424074
```

1. Utiliza el método de validación cruzada Monte Carlo

```
# Find histograms of MSE and MAE and R^2 using Shuffle Split (Monte Carlo)
kf = ShuffleSplit(n_splits=1000, test_size = 0.2)
mse_cv = []
mae_cv = []
r2_cv = []
for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]
    regr_cv = linear_model.LinearRegression()
    regr_cv.fit(x_train, y_train)
    # Test phase
    x_test = x[test_index, :]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)
    # Calculate MSE, MAE and R^2
    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)
    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)
    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)
fig, axs = plt.subplots(1, 3, sharey=True, tight_layout=True)
axs[0].hist(mse_cv, bins=25)
axs[0].set_xlabel("MSE")
axs[1].hist(mae_cv, bins=20)
axs[1].set_xlabel("MAE")
axs[2].hist(r2_cv, bins=20)
axs[2].set_xlabel("R^2")
plt.show()
```



1. Repite los pasos 1, 2 y 3 pero con este nuevo conjunto de datos

```
#Adding data
```

```
x = np.array([df['M'],df['W'],df['H'],df['P'],df['P']**2,  
df['H']**3,df['W']*df['M'],df['M']*df['P']**2]).T
```

```
X = np.column_stack((np.ones(x.shape[0]), x))
```

```
y = np.array(df['VR'])
```

```
beta = fit_model(X,y)
```

```
print("Coefficients:", beta)
```

```
y_pred = predict(X,beta)
```

```
Coefficients: [ 1.81681346e+03  1.88589943e+01  1.61293851e+00 -  
5.73914844e+01  
8.98146706e+01 -3.32605234e+00  4.02718450e-03 -1.83973900e-01  
2.95810994e-02]
```

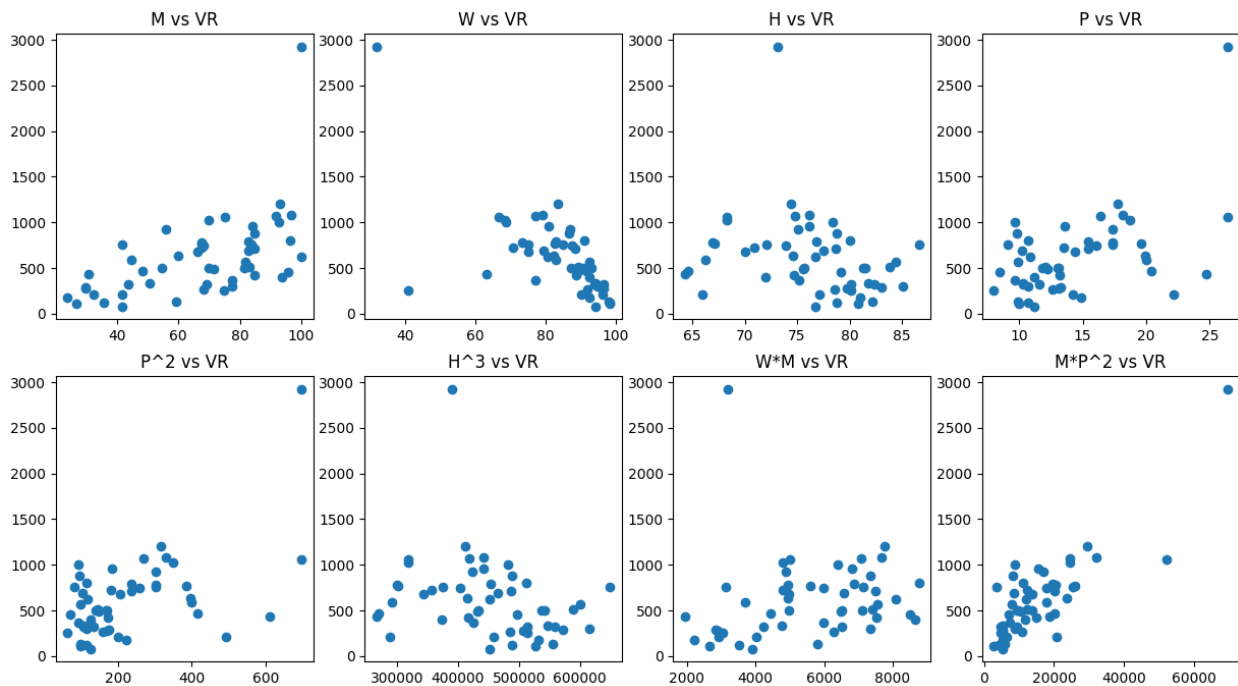
```
fig, ax = plt.subplots(2,4 , figsize=(15, 8))  
ax[0, 0].scatter(df['M'],df['VR'],label='Graph 1')  
ax[0, 0].set_title('M vs VR')  
ax[0, 1].scatter(df['W'],df['VR'],label='Graph 2')  
ax[0, 1].set_title('W vs VR')  
ax[0, 2].scatter(df['H'],df['VR'],label='Graph 3')  
ax[0, 2].set_title('H vs VR')
```

```

ax[0, 3].scatter(df['P'],df['VR'],label='Graph 4')
ax[0, 3].set_title('P vs VR')
ax[1,0].scatter(x[:,4],df['VR'],label='Graph 5')
ax[1, 0].set_title('P^2 vs VR')
ax[1,1].scatter(x[:,5],df['VR'],label='Graph 6')
ax[1, 1].set_title('H^3 vs VR')
ax[1,2].scatter(x[:,6],df['VR'],label='Graph 7')
ax[1, 2].set_title('W*M vs VR')
ax[1,3].scatter(x[:,7],df['VR'],label='Graph 8')
ax[1, 3].set_title('M*P^2 vs VR')

Text(0.5, 1.0, 'M*P^2 vs VR')

```



Evalúa con validación cruzada de k-pliegues tu modelo, calculando los valores de R2, MSE y MAE.

```

# Evaluate linear regression model using k-fold cross-validation

n_folds = 5
kf = KFold(n_splits=n_folds, shuffle = True)
mse_cv = []
mae_cv = []
r2_cv = []
for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

```

```

# print('a', x_train)
regr_cv = linear_model.LinearRegression()
regr_cv.fit(x_train, y_train)
# Test phase
x_test = x[test_index]
y_test = y[test_index]
y_pred = regr_cv.predict(x_test)
# Calculate MSE, MAE and R^2
mse_i = mean_squared_error(y_test, y_pred)
print('mse = ', mse_i)
mse_cv.append(mse_i)
mae_i = mean_absolute_error(y_test, y_pred)
print('mae = ', mae_i)
mae_cv.append(mae_i)
r2_i = r2_score(y_test, y_pred)
print('r^2= ', r2_i)
r2_cv.append(r2_i)
print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), ' R^2:',
      np.average(r2_cv))

mse = 110850.73244367399
mae = 221.80828204124353
r^2= -0.26495198627366157
mse = 37663.3804865193
mae = 140.334316189304
r^2= 0.23336222751751945
mse = 46165.120073654754
mae = 163.8010904779152
r^2= 0.5497168722482373
mse = 160636.6868088749
mae = 226.4788047302747
r^2= 0.7256749688378765
mse = 27942.953393637268
mae = 140.53140764633142
r^2= 0.7406595239532112
MSE: 76651.77464127204 MAE: 178.59078021701376 R^2:
0.3968923212566366

```

1. Implementa regresión Ridge con descenso de gradiente, y genera el gráfico de Ridge para el conjunto de datos original

```

import matplotlib.pyplot as plt
import numpy as np

# Generate synthetic data for demonstration
x = np.array([df['M'], df['W'], df['H'], df['P']]).T
X = np.column_stack((np.ones(x.shape[0]), x))
y = np.array(df['VR'])

n_features= 4

```

```

# Regularization parameter
alphas = np.logspace(-6, 6, 200) # Range of alpha values
coefs_path = []

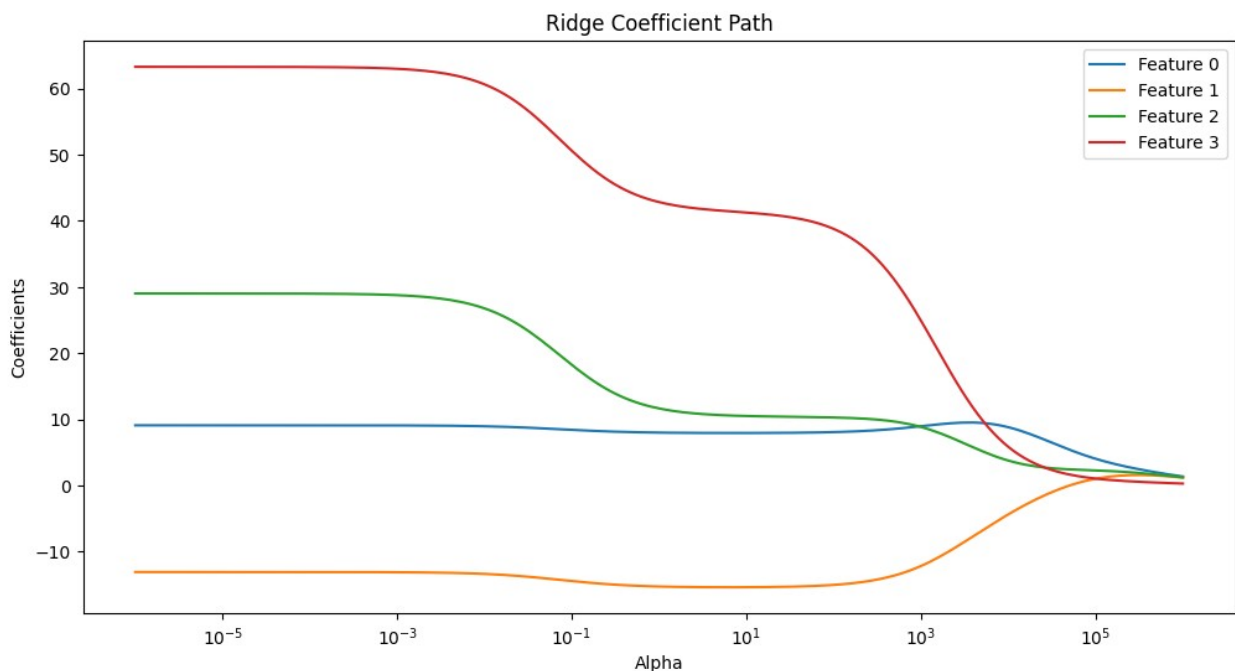
for a in alphas:
    XTX = np.dot(X.T, X)
    identity_matrix = np.identity(n_features + 1) # Identity matrix
    XTX_reg = XTX + a * identity_matrix
    XTX_inv_reg = np.linalg.inv(XTX_reg)
    coef = np.dot(np.dot(XTX_inv_reg, X.T), y)
    coefs_path.append(coef)

coefs_path = np.array(coefs_path) # Convert list of coefs to array

# Plot Ridge Coefficient Path (Ridge Graph)
plt.figure(figsize=(12, 6))
for feature_index in range(n_features):
    plt.plot(alphas, coefs_path[:, feature_index + 1], label=f'Feature {feature_index + 1}')

plt.xscale('log')
plt.xlabel('Alpha')
plt.ylabel('Coefficients')
plt.title('Ridge Coefficient Path')
plt.legend()
plt.axis('tight')
plt.show()

```



1. Utiliza una librería para generar el gráfico de Lasso

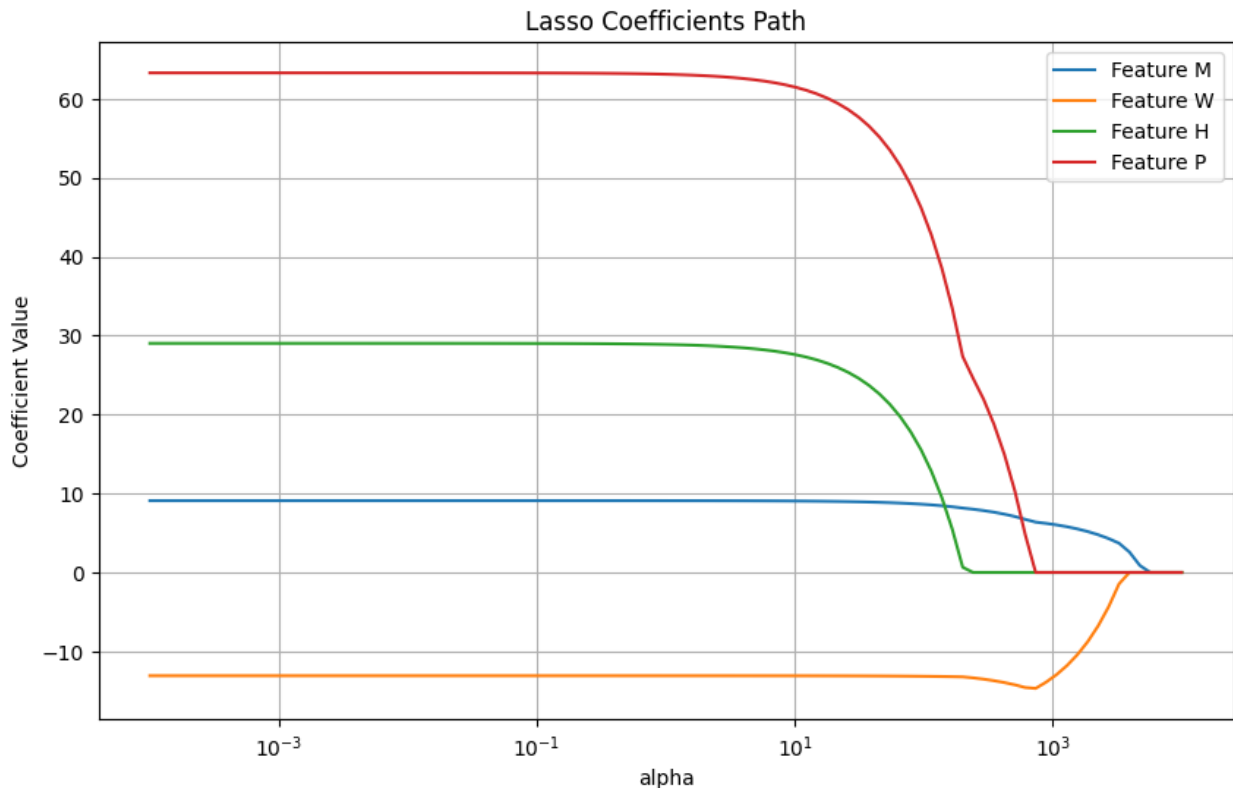
```
# Load the diabetes dataset
data = load_diabetes()
X = np.array([df['M'],df['W'],df['H'],df['P']]).T
y = np.array(df['VR'])
featuresDict = ['M', 'W', 'H', 'P']

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Define a range of alpha values
alphas = np.logspace(-4, 4, 100)

coefs = []
for alpha in alphas:
    lasso = Lasso(alpha=alpha)
    lasso.fit(X, y)
    coefs.append(lasso.coef_)

plt.figure(figsize=(10, 6))
for i in range(4):
    plt.plot(alphas, np.array(coefs)[: , i], label=f'Feature
{featuresDict[i]}')
plt.xscale('log')
plt.xlabel('alpha')
plt.ylabel('Coefficient Value')
plt.title('Lasso Coefficients Path')
plt.legend()
plt.grid(True)
plt.show()
```



Segun la grafica de Lasso, las variables mas importantes son el porcentaje de area metropolitana y el porcentaje de gente blanca

8. Conclusión problema 1

Considero que este modelo no es efectivo por que las r^2 obtenidas son inconsistentes, se puede conseguir tanto .80 como -.12 utilizando el modelo de regresion lineal, esto lo pude observar aplicando metodos de validación cruzada y Monte Carlo.

Cuando probe escalar las variables obtuve peores resultados, sin embargo esto se puede deber a las combinaciones que hice para este nuevo modelo (cuadrático y cubico), por otra parte se le agrega complejidad al modelo por que son mas variables y esto puede causar un peor resultado.

Para el modelo Lasso las variables en orden mas importantes son 'M', 'W', 'P' y la menos importante 'H'. Para el metodo de Ridge el resultado es parecido, su variable mas importante es 'W', posteriormente siguen 'M', 'H' y finalmente 'P'

Algo que me parecio interesante del modelo es que para cada variable hay un dato que esta muy diferente a los otros, por lo que es posible que esta medición afecte al modelo de una manera negativa ya que solo tenemos 50 mediciones.

Ejercicio 2

```
# Leer y limpiar datos
import pandas as pd
```

```

import numpy as np
df =
pd.read_csv('/home/alanv/Documents/7/omar/life_expectancy_data.csv')
df.drop(['Country', 'Year', 'percentage expenditure', 'Status', 'under-
five deaths ', ' HIV/AIDS', ' thinness 5-9 years'], inplace=True, axis=1)
df.dropna(inplace=True)
x = df.iloc[:, 1:].values
y = np.array(df['Life expectancy '])

```

1. Evalúa con validación cruzada un modelo de regresión lineal

#Limpiar datos

```

# Evalúa con validación cruzada un modelo de regresión
# lineal para las variables asignadas según tu matrícula utilizando
alguna librería o framework.
n_folds = 5
kf = KFold(n_splits=n_folds, shuffle = True)
mse_cv = []
mae_cv = []
r2_cv = []
for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]
    #print('a', x_train)
    regr_cv = linear_model.LinearRegression()
    regr_cv.fit(x_train, y_train)
    # Test phase
    x_test = x[test_index]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)
    # Calculate MSE, MAE and R^2
    mse_i = mean_squared_error(y_test, y_pred)
    print('mse = ', mse_i)
    mse_cv.append(mse_i)
    mae_i = mean_absolute_error(y_test, y_pred)
    print('mae = ', mae_i)
    mae_cv.append(mae_i)
    r2_i = r2_score(y_test, y_pred)
    print('r^2= ', r2_i)
    r2_cv.append(r2_i)
print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), ' R^2:',
np.average(r2_cv))

```

```

mse = 15.10303403591862
mae = 2.934046745120259
r^2= 0.8172720352424445
mse = 23.330594055730263

```

```

mae = 3.398286733478169
r^2= 0.6523871427021171
mse = 18.64093940094074
mae = 3.1328358770298963
r^2= 0.7574882076496504
mse = 15.849844270593602
mae = 2.9181180471366464
r^2= 0.7818505024823675
mse = 20.776120884679777
mae = 3.222722622041928
r^2= 0.7618997744999849
MSE: 18.740106529572603 MAE: 3.12120200496138 R^2: 0.754179532515313

```

1. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada

```

from sklearn.feature_selection import SelectKBest, r_regression
# Encuentra el número óptimo de predictores para el modelo utilizando
el método filter y validación cruzada.
# Una vez que tengas el número óptimo, muestra las características
seleccionadas.

x = df.iloc[:, 1:].values
y = np.array(df['Life expectancy '])
# Find optimal number of features using cross-validation
n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14]
mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)
    mse_cv = []
    mae_cv = []
    r2_cv = []
    kf = KFold(n_splits=5, shuffle = True)
    for train_index, test_index in kf.split(x):
        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]
        fselection_cv = SelectKBest(r_regression, k = n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)
        regr_cv = linear_model.LinearRegression()
        regr_cv.fit(x_train, y_train)
        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)
        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

```

```

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)
    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)
mae = np.average(mae_cv)
mae_nfeat.append(mae)
r2 = np.average(r2_cv)
r2_nfeat.append(r2)
print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("k")
axs[0].set_ylabel("MSE")
axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("k")
axs[1].set_ylabel("MAE")
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("k")
axs[2].set_ylabel("r^2")
plt.show()


---- n features = 1
MSE: 37.455501069927934 MAE: 4.556635172563096 R^2:
0.5150696032689794
---- n features = 2
MSE: 31.985579827566436 MAE: 4.177656516390863 R^2:
0.5854208158978527
---- n features = 3
MSE: 30.849662545479173 MAE: 4.132256441739769 R^2:
0.6013449541730767
---- n features = 4
MSE: 30.311233300057392 MAE: 4.091280203922175 R^2:
0.6073164307289786
---- n features = 5
MSE: 29.116611784082103 MAE: 4.04731786542821 R^2:
0.6229342040148208
---- n features = 6
MSE: 28.94731153191759 MAE: 4.025042152520049 R^2:
0.6239198003031836
---- n features = 7
MSE: 28.855804145294115 MAE: 4.02985511603051 R^2:
0.6224075731287417
---- n features = 8
MSE: 29.187205970287984 MAE: 4.036296344297083 R^2:

```

```

0.6214181173523587
---- n features = 9
MSE: 29.049730788925995  MAE: 4.032209765256544  R^2:
0.6227929100846179
---- n features = 10
MSE: 28.874847878583967  MAE: 4.028654082903706  R^2:
0.6261660470493798
---- n features = 11
MSE: 29.2458393264855  MAE: 4.0421767782837135  R^2:
0.6196301610973016
---- n features = 12
MSE: 28.98070427098433  MAE: 4.031906220250775  R^2:
0.6246404539471639
---- n features = 13
MSE: 28.709312042084527  MAE: 4.01860445834808  R^2:
0.6290695841784141
---- n features = 14
MSE: 18.598838583467245  MAE: 3.1128203039730544  R^2:
0.7602251565817171

```



En este caso las mejores variables serían todas, ya que cada vez que se incrementa el número de k aumenta la R^2 y disminuye tanto el MSE y MAE.

1. Encuentra el número óptimo de predictores para el modelo utilizando el método wrapper y validación cruzada

```
from sklearn.feature_selection import SequentialFeatureSelector
# Repite el paso anterior pero con selección de características
secuencial (Wrapper)

x = df.iloc[:, 1:].values
y = np.array(df['Life expectancy '])

# Find optimal number of features using cross-validation
n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)
    mse_cv = []
    mae_cv = []
    r2_cv = []
    kf = KFold(n_splits=5, shuffle = True)
    for train_index, test_index in kf.split(x):
        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]
        regr_cv = linear_model.LinearRegression()
        fselection_cv =
SequentialFeatureSelector(regr_cv, n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)
        regr_cv.fit(x_train, y_train)
        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)
        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)
        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)
        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)
    mae = np.average(mae_cv)
    mae_nfeat.append(mae)
    r2 = np.average(r2_cv)
    r2_nfeat.append(r2)
```

```


    print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")
axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")
plt.show()
x

---- n features = 1
MSE: 37.52908470907968  MAE: 4.24407844332222  R^2: 0.5090916921848996
---- n features = 2
MSE: 23.827012094368595  MAE: 3.2734313147248173  R^2:
0.6911733035928085
---- n features = 3
MSE: 19.950330073310074  MAE: 3.150131218184021  R^2:
0.7394894380558766
---- n features = 4
MSE: 19.32043734253997  MAE: 3.1396132624105593  R^2:
0.7474851100525348
---- n features = 5
MSE: 19.32243700836055  MAE: 3.137301406985551  R^2: 0.749427686211815
---- n features = 6
MSE: 18.837063097719295  MAE: 3.1081652201738312  R^2:
0.756619795189165
---- n features = 7
MSE: 18.717101699768445  MAE: 3.118701243043469  R^2:
0.7569657511407852
---- n features = 8
MSE: 18.60585490417259  MAE: 3.115109462514129  R^2:
0.7592986654755267
---- n features = 9
MSE: 18.80767448155364  MAE: 3.1286464746852936  R^2:
0.7551898440569662
---- n features = 10
MSE: 18.622485248910404  MAE: 3.1191519679268933  R^2:
0.7598279084229537
---- n features = 11
MSE: 18.70827100996998  MAE: 3.1238393021019357  R^2:
0.7563111060013713
---- n features = 12
MSE: 18.840521146954032  MAE: 3.1367450877008847  R^2:
0.7542842479958694
---- n features = 13

```


MSE: 18.6060425181228 MAE: 3.1081600155751987 R²: 0.7569585812500703



```
array([[2.63e+02, 6.20e+01, 1.00e-02, ..., 1.72e+01, 4.79e-01,
        1.01e+01],
       [2.71e+02, 6.40e+01, 1.00e-02, ..., 1.75e+01, 4.76e-01,
        1.00e+01],
       [2.68e+02, 6.60e+01, 1.00e-02, ..., 1.77e+01, 4.70e-01,
        9.90e+00],
       ...,
       [7.30e+01, 2.50e+01, 4.43e+00, ..., 1.20e+00, 4.27e-01,
        1.00e+01],
       [6.86e+02, 2.50e+01, 1.72e+00, ..., 1.60e+00, 4.27e-01,
        9.80e+00],
       [6.65e+02, 2.40e+01, 1.68e+00, ..., 1.10e+01, 4.34e-01,
        9.80e+00]])
```

Viendo los resultados se pueden asumir conclusiones diferentes al metodo filter, por un lado se puede ver con 5 variables que ya no mejora mucho mas el modelo, sin embargo se puede ver un pico notorio para 10 variables, por eso es por lo que lo eligo, las variables son las siguientes:

```

regr = linear_model.LinearRegression()
fselection = SequentialFeatureSelector(regr, n_features_to_select =
10)
fselection.fit(x, y)
print("Selected features: ", fselection.get_feature_names_out())
df.columns[1:]

Selected features:  ['x0' 'x1' 'x2' 'x3' 'x5' 'x8' 'x9' 'x10' 'x12'
'x13']

Index(['Adult Mortality', 'infant deaths', 'Alcohol', 'Hepatitis B',
      'Measles ', ' BMI ', 'Polio', 'Total expenditure', 'Diphtheria
', 'GDP',
      'Population', ' thinness 1-19 years',
      'Income composition of resources', 'Schooling'],
      dtype='object')

```

1. Encuentra el número óptimo de predictores para el modelo utilizando el método filter wrapper

```

#Ahora con filter wrapper
from sklearn.feature_selection import RFE
# Find optimal number of features using cross-validation
x = df.iloc[:, 1:].values
y = np.array(df['Life expectancy '])
n_feats = [1,2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,14]
mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)
    mse_cv = []
    mae_cv = []
    r2_cv = []
    kf = KFold(n_splits=5, shuffle = True)
    for train_index, test_index in kf.split(x):
        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]
        regr_cv = linear_model.LinearRegression()
        fselection_cv = RFE(regr_cv,n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)
        regr_cv.fit(x_train, y_train)
        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)
        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

```

```

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)
    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)
mae = np.average(mae_cv)
mae_nfeat.append(mae)
r2 = np.average(r2_cv)
r2_nfeat.append(r2)
print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")
axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")
plt.show()


---- n features = 1
MSE: 37.335739236591394 MAE: 4.246104185145866 R^2:
0.5156841735621317
---- n features = 2
MSE: 32.10754487569488 MAE: 4.183422370875017 R^2: 0.584838932048305
---- n features = 3
MSE: 31.34842796470945 MAE: 4.157807356154661 R^2:
0.5913337251929941
---- n features = 4
MSE: 30.48964864079943 MAE: 4.138542225896471 R^2:
0.6043610378353089
---- n features = 5
MSE: 30.28944961842293 MAE: 4.12993475267375 R^2: 0.6072805711696848
---- n features = 6
MSE: 21.674591495508576 MAE: 3.3540860324400796 R^2:
0.719702119643586
---- n features = 7
MSE: 19.097913117128265 MAE: 3.1340157649604157 R^2:
0.7523853911421309
---- n features = 8
MSE: 18.913204199034976 MAE: 3.115664553577293 R^2:
0.7547690428039233
---- n features = 9
MSE: 18.98299041548599 MAE: 3.121196344669433 R^2:
0.7511612875568068

```

```

---- n features = 10
MSE: 18.887832207389536  MAE: 3.116204800187327  R^2:
0.7547755959769972
---- n features = 11
MSE: 19.11345646337707  MAE: 3.128843466688746  R^2:
0.7511888821446613
---- n features = 12
MSE: 18.623822895094968  MAE: 3.105386632739015  R^2:
0.7590811384941188
---- n features = 13
MSE: 18.60654287124803  MAE: 3.107300808170573  R^2:
0.7578261584386352
---- n features = 14
MSE: 18.726036255266365  MAE: 3.118314873272967  R^2:
0.7551884584390431

```



En este caso se nota mucho que con 7 variables es donde esta el mayor pico, estas 7 variables son:

```

regr = linear_model.LinearRegression()
fselection = RFE(regr, n_features_to_select = 7)
fselection.fit(x, y)

```

```
print("Selected features: ", fselection.get_feature_names_out())
df.columns[1:]

Selected features:  ['x0' 'x2' 'x5' 'x8' 'x11' 'x12' 'x13']

Index(['Adult Mortality', 'infant deaths', 'Alcohol', 'Hepatitis B',
      'Measles ', ' BMI ', 'Polio', 'Total expenditure', 'Diphtheria
', 'GDP',
      'Population', ' thinness 1-19 years',
      'Income composition of resources', 'Schooling'],
      dtype='object')
```

1. Impelentacion de los pasos anteriores pero con un modelo no lineal knn

Impelentacion de los pasos anteriores pero con un modelo no lineal como knn

```
import numpy as np
from sklearn.model_selection import KFold
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

x = df.iloc[:, 1:].values
y = np.array(df['Life expectancy '])

n_folds = 5
kf = KFold(n_splits=n_folds, shuffle=True)
mse_cv = []
mae_cv = []
r2_cv = []

for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    # Initialize kNN regressor
    k_neighbors = 5 # You can adjust this value
    regr_cv = KNeighborsRegressor(n_neighbors=k_neighbors)
    regr_cv.fit(x_train, y_train)

    # Test phase
    x_test = x[test_index]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    # Calculate MSE, MAE, and R^2
    mse_i = mean_squared_error(y_test, y_pred)
    print('mse = ', mse_i)
    mse_cv.append(mse_i)
```

```

mae_i = mean_absolute_error(y_test, y_pred)
print('mae = ', mae_i)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
print('r^2 = ', r2_i)
r2_cv.append(r2_i)

print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), '
R^2:', np.average(r2_cv))

```

```

mse = 84.63127999999999
mae = 7.004969696969696
r^2 = 0.01067748114694389
mse = 73.14774181818181
mae = 6.577151515151515
r^2 = 0.0060996746468604535
mse = 64.08419393939396
mae = 6.137333333333333
r^2 = 0.09190193479625741
mse = 77.90113212121211
mae = 6.807515151515152
r^2 = 0.05028438337446184
mse = 69.1899623100304
mae = 6.395866261398176
r^2 = 0.07048016205706742
MSE: 73.79086203776366 MAE: 6.584567191673574 R^2:
0.045888727204318205

```

```

from sklearn.feature_selection import SelectKBest, f_regression

```

```

# Find optimal number of features using cross-validation
n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14]
mse_nfeat = []
mae_nfeat = []
r2_nfeat = []

```

```

for n_feat in n_feats:
    print('---- n features =', n_feat)
    mse_cv = []
    mae_cv = []
    r2_cv = []
    kf = KFold(n_splits=5, shuffle=True)

    for train_index, test_index in kf.split(x):
        # Training phase

```

```

x_train = x[train_index, :]
y_train = y[train_index]

# Feature selection
fselection_cv = SelectKBest(k=n_feat)
fselection_cv.fit(x_train, y_train)
x_train = fselection_cv.transform(x_train)

# Initialize kNN regressor
k_neighbors = 5 # You can adjust this value
regr_cv = KNeighborsRegressor(n_neighbors=k_neighbors)
regr_cv.fit(x_train, y_train)

# Test phase
x_test = fselection_cv.transform(x[test_index, :])
y_test = y[test_index]
y_pred = regr_cv.predict(x_test)

mse_i = mean_squared_error(y_test, y_pred)
mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)


r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("k")
axs[0].set_ylabel("MSE")
axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("k")
axs[1].set_ylabel("MAE")
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("k")
axs[2].set_ylabel("r^2")
plt.show()

```

```
---- n features = 1
MSE: 26.60335088514322 MAE: 3.476102938196555 R^2:
0.6481571369252659
---- n features = 2
MSE: 10.95621727254306 MAE: 2.0750958828405635 R^2:
0.8571284728960661
---- n features = 3
MSE: 8.717474146449296 MAE: 1.9889184489269596 R^2:
0.8877170186211438
---- n features = 4
MSE: 21.78770286966934 MAE: 3.2432942801878974 R^2:
0.7163623103256391
---- n features = 5
MSE: 19.748631909367226 MAE: 3.082095311780418 R^2:
0.7433982385470829
---- n features = 6
MSE: 18.80149510840932 MAE: 3.0108329372754907 R^2:
0.7572270050543397
---- n features = 7
MSE: 18.81438288809063 MAE: 3.00252950170397 R^2: 0.7570278143985534
---- n features = 8
MSE: 19.180937300543427 MAE: 3.0614543243990058 R^2:
0.7519885175418439
---- n features = 9
MSE: 19.428318524085846 MAE: 3.0891391360412634 R^2:
0.746063557183105
---- n features = 10
MSE: 18.684388504375057 MAE: 3.035140793957815 R^2:
0.7551657001049124
---- n features = 11
MSE: 20.436363078566824 MAE: 3.1766830247766413 R^2:
0.7343804903452004
---- n features = 12
MSE: 20.856067803628996 MAE: 3.2016603481624757 R^2:
0.7285348749664812
---- n features = 13
MSE: 38.74065444155845 MAE: 4.420861121856867 R^2:
0.4902917859904778
---- n features = 14
MSE: 72.29103814460719 MAE: 6.5121108961960035 R^2:
0.06349132400739557
```

```
fselection = SelectKBest(f_regression, k=2)
fselection.fit(x, y)
selected_features = fselection.get_support()

# Get the column names from the dataframe
selected_column_names = df.columns[1:][selected_features]

print("Mejores variables: ", selected_column_names)

Mejores variables: Index(['Income composition of resources',
'Schooling'], dtype='object')

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14]
mse_nfeat = []
mae_nfeat = []
r2_nfeat = []

for n_feat in n_feats:
    print('---- n features =', n_feat)
    mse_cv = []
    mae_cv = []
    r2_cv = []
    kf = KFold(n_splits=5, shuffle=True)
```

```

for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    # Initialize kNN regressor
    k_neighbors = 5
    regr_cv = KNeighborsRegressor(n_neighbors=k_neighbors)
    regr_cv.fit(x_train, y_train)

    # Feature selection
    fselection_cv = SelectKBest(score_func=f_regression, k=n_feat)
    fselection_cv.fit(x_train, y_train)
    x_train_selected = fselection_cv.transform(x_train)

    # Test phase
    x_test = x[test_index, :]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

```


```

fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")
axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")

```

```
axs[2].set_ylabel("r^2")  
plt.show()
```

```
---- n features = 1  
MSE: 72.1459044008474 MAE: 6.561898646034817 R^2:  
0.06301027562261421  
---- n features = 2  
MSE: 70.09538102385557 MAE: 6.460312572533849 R^2:  
0.09325686236391909  
---- n features = 3  
  
MSE: 73.94816353688864 MAE: 6.632416542322924 R^2:  
0.043744509673168584  
---- n features = 4  
MSE: 73.42714732835958 MAE: 6.627289085382704 R^2:  
0.049556437747119333  
---- n features = 5  
MSE: 76.55777862319242 MAE: 6.747271990420927 R^2:  
0.006079578035025834  
---- n features = 6  
MSE: 73.69114042774247 MAE: 6.622709809339597 R^2:  
0.04374072537584182  
---- n features = 7  
MSE: 73.91578734457032 MAE: 6.665622768720641 R^2:  
0.042749586999220934  
---- n features = 8  
MSE: 72.01286015142304 MAE: 6.523809781707655 R^2:  
0.06826496734287804  
---- n features = 9  
MSE: 73.40219227852998 MAE: 6.621072708851432 R^2:  
0.04750374080728437  
---- n features = 10  
MSE: 73.43466185281386 MAE: 6.627512462006078 R^2:  
0.0464887221339507  
---- n features = 11  
MSE: 70.81025091719627 MAE: 6.496759988947223 R^2:  
0.0795765633138241  
---- n features = 12  
MSE: 71.69684510785666 MAE: 6.510692640692641 R^2:  
0.0702696386430595  
---- n features = 13  
MSE: 72.99370889564338 MAE: 6.52449704338215 R^2:  
0.05156235329880181  
---- n features = 14  
MSE: 72.12953694759142 MAE: 6.572603628995118 R^2:  
0.0631677857957234
```



El mejor numero de features son 8 aun asi el modelo es muy malo

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import KFold
from sklearn.feature_selection import SelectKBest,
mutual_info_regression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14]
mse_nfeat = []
mae_nfeat = []
r2_nfeat = []

for n_feat in n_feats:
    print('---- n features =', n_feat)
    mse_cv = []
    mae_cv = []
    r2_cv = []
    kf = KFold(n_splits=5, shuffle=True)
```

```

for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    # Initialize kNN regressor
    k_neighbors = 5
    regr_cv = KNeighborsRegressor(n_neighbors=k_neighbors)
    regr_cv.fit(x_train, y_train)

    # Feature selection using SelectKBest and mutual information
    fselection_cv = SelectKBest(score_func=mutual_info_regression,
k=n_feat)
    fselection_cv.fit(x_train, y_train)
    x_train_selected = fselection_cv.transform(x_train)

    # Test phase
    x_test = x[test_index, :]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)


print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")
axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

```

```
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")
plt.show()
```

```
---- n features = 1
MSE: 70.37027591489363 MAE: 6.457421645021645 R^2:
0.08821455391669183
---- n features = 2
MSE: 74.65363425329281 MAE: 6.641304006631667 R^2:
0.028649153279945085
---- n features = 3
MSE: 72.79552082601087 MAE: 6.592485106382978 R^2:
0.05752468633117036
---- n features = 4
MSE: 71.64643229547758 MAE: 6.507330533296491 R^2:
0.06572899738738705
---- n features = 5
MSE: 70.58776955991526 MAE: 6.50327331675417 R^2:
0.08231295986795042
---- n features = 6
MSE: 74.14073064382427 MAE: 6.648587381412914 R^2:
0.03628026215519513
---- n features = 7
MSE: 72.44565596094685 MAE: 6.557483724785851 R^2:
0.06022655817811249
---- n features = 8
MSE: 72.65026522612139 MAE: 6.5717697706548766 R^2:
0.05512199018904338
---- n features = 9
MSE: 70.94178979976053 MAE: 6.492715851524362 R^2:
0.08124294203736841
---- n features = 10
MSE: 72.27029739780787 MAE: 6.550156323109515 R^2:
0.06554098407319582
---- n features = 11
MSE: 73.43536242940039 MAE: 6.57519959473151 R^2:
0.04667639562089121
---- n features = 12
MSE: 72.93447265064013 MAE: 6.575659427097724 R^2:
0.0507435021655299
---- n features = 13
MSE: 72.30043184305057 MAE: 6.580009321175278 R^2:
0.06102980499226749
---- n features = 14
MSE: 72.71140810684352 MAE: 6.571369512756746 R^2:
0.05397069866930666
```



1. Agregue la variables "Status" y utiliza un árbol de decisión para generar un modelo de regresión

```
dfTree =
pd.read_csv('/home/alanv/Documents/7/omar/life_expectancy_data.csv')
dfTree.drop(['Country', 'Year', 'percentage expenditure', 'under-five
deaths ', ' HIV/AIDS', ' thinness 5-9 years'], inplace=True, axis=1)
dfTree.dropna(inplace=True)
x = dfTree.iloc[:, 1:].values
y = np.array(dfTree['Life expectancy '])

import numpy as np
from sklearn.model_selection import KFold
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

n_folds = 5
kf = KFold(n_splits=n_folds, shuffle=True)
mse_cv = []
mae_cv = []
```

```

r2_cv = []

# Hyperparameters to adjust
max_depth = 10
min_samples_split = 2
min_samples_leaf = 1

for train_index, test_index in kf.split(x):
    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    # Initialize Decision Tree Regressor with adjusted hyperparameters
    regr_cv = DecisionTreeRegressor(
        max_depth=max_depth,
        min_samples_split=min_samples_split,
        min_samples_leaf=min_samples_leaf
    )
    regr_cv.fit(x_train, y_train)

    # Test phase
    x_test = x[test_index]
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    # Calculate MSE, MAE and R^2
    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), ' R^2:', np.average(r2_cv))

MSE: 0.007639421571336476 MAE: 0.028985907709313852 R^2: 0.9999015930259334

```

7. Conclusión problema 2

Yo considero que sí es un modelo aceptable, a diferencia del problema anterior, ya que el valor de r^2 es considerablemente alto y los errores son aceptables, además de que los resultados son estables y no hay mucha variabilidad.

Todos los métodos lineales funcionan correctamente, dan un valor de r^2 parecido aunque con diferente número de variables. Por otra parte, el único modelo no lineal que me dio un buen resultado fue con el método filter, y probablemente sea porque este modelo no tiene en cuenta

las interacciones entre variables; este mismo fue el modelo más sobresaliente, con una suma de errores cuadrados de 2, algo que destacar es que el número de variables recomendadas son solamente 2 o 3. Al agregar la variable Status, mejora mucho el modelo, con unos errores muy bajos y una r cuadrada de 0.9999.

Me pareció interesante que para cada método tiene un número de variables óptimas diferentes al resto. Hay algunos que recomiendan usar todas las características, mientras que otros solo recomiendan 2. Para los métodos lineales el valor de r cuadrada es prácticamente el mismo pero cuando tiene el mejor número de variables, filter fue con el máximo número de variables, wrapper con 8 y filter-wrapper con 7, el método de filter se debe a que como deprecia la interacción entre variables el que tenga más no hace más complejo el modelo, caso contrario para los otros métodos.